

Exception handling for the BLAS and LAPACK

Wesley da Silva Pereira, University of Colorado Denver

September 14, 2022

XBLAS meeting

Examples of bad exception handling

- Ariane 501, 1996. “[An] exception was caused during execution of a data conversion from 64-bit floating point to 16-bit signed integer value.” ²
- Roborace, 2020. “the system somehow managed to produce a NaN (not a number) value and all verification logic was designed to work only with numbers.” ³



²<https://www-users.cse.umn.edu/~arnold/disasters/ariane.html>

³<https://www.reddit.com/r/formula1/comments/jk9jrg/comment/gai2951>

Agenda

- ① Inconsistencies in current BLAS and LAPACK libraries
- ② Proposed exception handling and preliminary tests
- ③ Concluding remarks

Goals of this talk:

Show inconsistencies in current BLAS and LAPACK libraries

Propose a solution

Seek feedback

Inconsistencies in the BLAS

l_xAMAX: Return the index of the largest entry.

- Inconsistency #1:

`isamax([0, NaN, 2]) = 3,`

`isamax([NaN, 0, 2]) = 1,`

- Inconsistency #2:

`icamax([OV+i*OV, Inf+i*0]) = 1,`

`icamax([.6*OV+i*.6*OV, .7*OV+i*.7*OV]) = 1,`

where OV is the overflow threshold.

- Proposed consistent fix:

**Return index to the first NaN,
else to the first Inf,
else to the first largest finite entry.**

Inconsistencies in the BLAS

xGER: $A := A + \alpha xy^T$.

- Inconsistency: Depending on the value of α or y_j , Infs and NaNs do not propagate.
- Current Standard BLAS:
 - If $\alpha = 0$, return A .
 - If $y_j = 0$ then $A_{ij} := A_{ij}$.
 - There is no check for $x_i = 0$.
- Proposed consistent fix:

Keep check for $\alpha = 0$ but not for $y_j = 0$ (or $x_i = 0$).

Inconsistencies in LAPACK

xGETF2 + xGETRS: Solve $Ax = b$ for x .

- Example: $A = \begin{bmatrix} 1 & 0 \\ NaN & 2 \end{bmatrix}$ and $b = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$.
- Combination of BLAS inconsistencies:
 - IxAMAX chooses $A_{11} = 1$, not $A_{21} = NaN$, as pivot.

$$A = \begin{bmatrix} 1 & 0 \\ NaN & 2 \end{bmatrix}$$

Inconsistencies in LAPACK

xGETF2 + xGETRS: Solve $Ax = b$ for x .

- Example: $A = \begin{bmatrix} 1 & 0 \\ NaN & 2 \end{bmatrix}$ and $b = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$.
- Combination of BLAS inconsistencies:
 - IxAMAX chooses $A_{11} = 1$, not $A_{21} = NaN$, as pivot.
 - xGER finds 0 in A_{12} and do not use the NaN.

$$A_{22} := A_{22} - NaN \cdot 0$$

Inconsistencies in LAPACK

xGETF2 + xGETRS: Solve $Ax = b$ for x .

- Example: $A = \begin{bmatrix} 1 & 0 \\ NaN & 2 \end{bmatrix}$ and $b = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$.
- Combination of BLAS inconsistencies:
 - IxAMAX chooses $A_{11} = 1$, not $A_{21} = NaN$, as pivot.
 - xGER finds 0 in A_{12} and do not use the NaN.
 - xGETRS sees the LU factorization:

$$L = \begin{bmatrix} 1 & 0 \\ NaN & 1 \end{bmatrix}, \quad U = \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

Inconsistencies in LAPACK

xGETF2 + xGETRS: Solve $Ax = b$ for x .

- Example: $A = \begin{bmatrix} 1 & 0 \\ NaN & 2 \end{bmatrix}$ and $b = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$.
- Combination of BLAS inconsistencies:
 - IxAMAX chooses $A_{11} = 1$, not $A_{21} = NaN$, as pivot.
 - xGER finds 0 in A_{12} and do not use the NaN.
 - xGETRS sees the LU factorization:

$$L = \begin{bmatrix} 1 & 0 \\ NaN & 1 \end{bmatrix}, \quad U = \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

- The first call for xTRSV solves $Ly = b$ and finds $y = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$. In the second iteration, $y_2 := 1 - NaN \cdot y_1$ and $y_1 = 0$.

Inconsistencies in LAPACK

xGETF2 + xGETRS: Solve $Ax = b$ for x .

- Example: $A = \begin{bmatrix} 1 & 0 \\ NaN & 2 \end{bmatrix}$ and $b = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$.
- Combination of BLAS inconsistencies:
 - IxAMAX chooses $A_{11} = 1$, not $A_{21} = NaN$, as pivot.
 - xGER finds 0 in A_{12} and do not use the NaN.
 - xGETRS sees the LU factorization:

$$L = \begin{bmatrix} 1 & 0 \\ NaN & 1 \end{bmatrix}, \quad U = \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$$

- The first call for xTRSV solves $Ly = b$ and finds $y = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$.
- The final solution is $x = \begin{bmatrix} 0 \\ 0.5 \end{bmatrix}$
- **NaN does not propagate!**

Examples of unavoidable inconsistencies

IEEE 754-2019 defines min and max that propagate NaNs.

Compilers need time to adopt those rules. ⁴

```
1 min( qNaN, 1.0 ) = qNaN ! gfortran 12.2
2 min( 1.0, qNaN ) = 1.0 ! gfortran 12.2
```

Absolute value, product, and quotient of complex numbers:

- They are not specified by IEEE 754-2019.
- In C99, C++ and Fortran (starting in 2008): computations “without undue overflow or underflow”.
- There is space for inconsistencies, e.g., ⁵

```
1 cmplx(Inf,0.0)*cmplx(Inf,Inf) = cmplx(Inf,Inf) ! ifort 2021
2 cmplx(Inf,0.0)*cmplx(Inf,Inf) = cmplx(NaN,NaN) ! gfortran 12
3
4 abs(cmplx(0.0,subNormal)) = subNormal ! ifort 2021 (-00)
5 abs(cmplx(0.0,subNormal)) = 0.0       ! ifort 2021 (-03)
```

⁴<https://www.godbolt.org/z/1sfceYbao>

⁵<https://www.godbolt.org/z/nvWWeToY6>

Proposed exception handling

If NaNs or Infs are inputs or created internally,

- ① The program will still terminate.
- ② Moreover, either:
 - NaNs and Infs propagate to the output.
 - NaNs and Infs are dealt with internally.
 - NaNs and Infs **do not propagate** to the output **in special cases**, e.g., $C := 0 \cdot AB + 0 \cdot C$ in GEMM.
- ③ LAPACK routines will report (using INFO and more).

Current design ⁶ takes into account:

- Several user requests.
- Different stakeholders.
- Relevant xSDK Community Policies for Library Development.

⁶<https://www.arxiv.org/abs/2207.09281>

The EC (Error Checking) interfaces

- Example:
`SGESV_EC(N, NRHS, A, LDA, IPIV, B, LDB, INFO,
FLAG_REPORT, INFO_ARRAY, CONTEXT)`
- 3 more arguments at end:
 - `FLAG_REPORT(1:2)`
`FLAG_REPORT(1)`: WHAT to report.
`FLAG_REPORT(2)`: HOW to report.
 - `INFO_ARRAY(:)` - can report `INFO` + details on Infs and NaNs on input, and internal calls.
 - `CONTEXT` - opaque argument. Can report, for example, exceptions on different threads.
- Legacy interface will be maintained as a wrapper.

Some early testBLAS⁸ results

Expected output from IxAMAX: Return index to the first NaN, else to the first Inf, else to the first largest finite entry.

| BLAS library | I{S,D}AMAX | I{C,Z}AMAX |
|-------------------------|--------------------|------------|
| Apple Accelerate 12.2.1 | DNFPS ⁷ | DNFPS |
| BLIS 0.9.0 | Pass | DNFPS |
| IBM ESSL 6.3.0 | DNFPS | DNFPS |
| Intel MKL 2022.1.0 | Pass | DNFPS |
| LAPACK 3.9.1 | DNFPS | DNFPS |
| LAPACK 3.10.1 | DNFPS | DNFPS |
| LAPACK 3.11-beta | Pass | Pass |
| OpenBLAS 0.3.8 | DNFPS | DNFPS |

⁷DNFPS: Does not follow proposed standard.

⁷Apple Accelerate on macOS Monterey 12.2.1 using Apple clang version 13.1.6. IBM ESSL on Summit node using GNU compiler v9.1.0. Others on Ubuntu 20.04.4 LTS using GNU compiler v9.4.0.

⁸<https://www.github.com/tlapack/testBLAS>

Some early testBLAS¹⁰ results

Expected output from xNRM2: NaN if in input, else Inf if in input, else “accurate answer” (possibly Inf).

| BLAS library | {S,D}NRM2 (finite input) | {S,D}NRM2 |
|-------------------------|--------------------------|-----------|
| Apple Accelerate 12.2.1 | DNFPS ⁹ | DNFPS |
| BLIS 0.9.0 | DNFPS | DNFPS |
| IBM ESSL 6.3.0 | Pass | Pass |
| Intel MKL 2022.1.0 | Pass | Pass |
| LAPACK 3.9.1 | DNFPS | DNFPS |
| LAPACK 3.10.1 | Pass | Pass |
| LAPACK 3.11-beta | Pass | Pass |
| OpenBLAS 0.3.8 | Pass | Pass |

⁹DNFPS: Does not follow proposed standard.

⁹Apple Accelerate on macOS Monterey 12.2.1 using Apple clang version 13.1.6. IBM ESSL on Summit node using GNU compiler v9.1.0. Others on Ubuntu 20.04.4 LTS using GNU compiler v9.4.0.

¹⁰<https://www.github.com/tlapack/testBLAS>

Proposed tasks

- 1 During installation of LAPACK, check `abs`, multiplication and division of complex, and `min` and `max` behave as expected.
(partially done in LAPACK 3.10.1)
- 2 Modify LAPACKE drivers, so that they return an error flag if input has NaNs or Infs.
- 3 Modify LAPACK, so that routines that compute norms of complex matrices signal NaNs and Infs on input.
- 4 Fix Reference BLAS and implement test code.
- 5 Validate the EC interfaces for a few LAPACK routines.
- 6 Design more general test code.
- 7 Implement the EC interfaces to the rest of LAPACK.

Thank you!

Questions?

Proposed Exception Handling for the BLAS and LAPACK:

<https://www.arxiv.org/abs/2207.09281>

Funded by:

NSF under the project Basic ALgebra Libraries for Sustainable Technology with Interdisciplinary Collaboration (BALLISTIC).



National
Science
Foundation

Meanings of WHAT and HOW

- $\text{WHAT} \leq -1$: check nothing.
 - $\text{WHAT} = 0$: legacy INFO checking.
 - $\text{WHAT} = 1$: also, check input/output args for Infs/NaNs.
 - $\text{WHAT} \geq 2$: do the above throughout call tree.
-
- $\text{HOW} \leq 0$: only use INFO.
 - $\text{HOW} = 1$: also use INFO_ARRAY.
 - $\text{HOW} = 2$: also, if $\text{INFO} \neq 0$, call
 `REPORT_EXCEPTIONS (CONTEXT, ROUTINE_NAME,
 INFO_ARRAY)`
 - $\text{HOW} = 3$: do the above throughout call tree.
 - $\text{HOW} \geq 4$: call
 `GET_FLAGS_TO_REPORT(CONTEXT, FLAG_REPORT)`

Tests corner cases and Inf/NaN propagation on BLAS packages.

- Uses the C++ BLAS interface of $\langle T \rangle$ LAPACK¹¹.
- Uses BLAS++¹² wrappers as main interface to vendor BLAS.
- Currently:
 - Test corner cases on BLAS routines.
 - Specific Inf and NaN propagation tests for `iamax`, `nrm2`, `trsv` and `trsm`.

¹¹<https://github.com/tlapack/tlapack>

¹²<https://bitbucket.org/icl/blaspp>

¹³<https://www.github.com/tlapack/testBLAS> 